

Computational Robotics

Lab 1 - Markov Decision Processes

Andrew Choi

October 2019

1 Introduction

This goal of this lab is to explore finite Markov Decision Processes (MDPs) to control a simple discretized robot. To do so, a gridworld is constructed with well-defined boundaries and rewards. In this gridworld, we develop and implement a model for robot behavior and use it to accomplish a prescribed task.

2 Mathematical Formulation

Before delving into the details of the experiment, the finite Markov Decision Process as well as the methods employed (value iteration and policy iteration) are described mathematically.

Markov decision processes (MDPs from hereon) are a classical formalization of sequential decision making. Driven by rewards, these decisions are made not just influenced by immediate rewards, but also by possible future rewards.

The finite MDP can be formulated as a set of states, actions, and rewards.

3 Experimental Setup

3.1 Robot Description

For the experiment, a simple robot is considered that will reside in the gridworld and attempt to reach a prescribed goal state. The robot will occupy one grid in the gridworld at any given time and can face any of the twelve headings identified by the hours on a clock $h \in \{0 \dots 11\}$. Each action is characterized by moving forward, backward, or staying still. This is followed by a rotation of -1, 0, or 1. From this it can be seen that the action of the robot can be described as a tuple

$$(a, b),$$

where a is the movement value and b is the rotation value. The following restrictions are also applied to the robot's actions:

1. Attempting to move off the grid results in no linear movement, but the rotation portion will still happen.
2. Aside from restriction 1, the robot can only rotate if it also moves forwards or backwards.
3. If the robot chooses to move "forwards" or "backwards" a pre-rotation error will occur with probability p_e
 - (a) With probability p_e each, robot will first rotate by +1 or -1 before it moves. No pre-rotate with probability $1 - p_e$.
 - (b) Choosing to stay still will not incur an error rotation.

Then, it can be seen that the robot is defined by the following discrete action space.

$$(1, 0), (1, 1), (1, -1), (0, 0), (-1, 0), (-1, 1), (-1, -1) N_e = 7$$




Figure 1: Gridworld with set Rewards

3.2 Created Gridworld and Rewards

To allow for trajectory simulation, an 8x8 gridworld is constructed. Provided for us are the defined rewards for the states of the gridworld as shown in figure 1.

As the rewards for each state are independent of heading angle, they can be represented graphically on a 2D grid image. As shown above, the boundary states (red and marked X) have a reward of -100. The lane markers

3.3 Benchmark

To be able to properly analyze the benefits provided by following an optimal policy, we first must set a benchmark. This benchmark will be the trajectory generated by a hand-engineered initial policy π_0 . This policy crudely gets to the goal state by simply prioritizing closing the x distance and then the y distance without any regard for rewards (*code in utils.py*). The plotted trajectory can be seen in figure 2.

3.4 Scenarios

With the following setup, we will explore 3 different scenarios:

1. Reward at goal state independent of heading, $p_e = 0\%$
2. Reward at goal state independent of heading, $p_e = 25\%$
3. Reward at goal state +1 only when robot is pointing down $h = 6$, $p_e = 0\%$
4. Reward at goal state +1 only when robot is pointing down $h = 6$, $p_e = 25\%$

4 Experimental Results

As discussed in the section 1, we employ dynamic programming to find the optimal value function of this MDP. First, we employ policy iteration.

4.1 Policy Iteration

First, we define our initial policy π_0 as the initial policy discussed in section 2.2, initial state s_0 as $x = 1, y = 6, h = 6$ (i.e top left corner, pointing down), and error probability $p_e = 0$. Using a discount factor $\gamma = 0.9$, we get the following trajectory.

4.2 Value Iteration

For value iteration, we first set the initial value function $V(s) = 0 \forall s \in S$.

4.3 Comparisons

First,

5 Conclusion